

Housing Price Regression

Author - Kainshk Karam

=====

STEP1: Import required libraries

```
In [2]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from warnings import filterwarnings
filterwarnings('ignore')
from sklearn.preprocessing import LabelEncoder
```

STEP2: Import file

```
In [3]: df=pd.read_csv("Housing.csv")
df.head(2)
```

```
Out[3]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea
0	13300000	7420	4	2	3	yes	no	no	no	yes	2	
1	12250000	8960	4	4	4	yes	no	no	no	yes	3	

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype  
---  -
0   price                545 non-null   int64  
1   area                 545 non-null   int64  
2   bedrooms             545 non-null   int64  
3   bathrooms             545 non-null   int64  
4   stories               545 non-null   int64  
5   mainroad             545 non-null   object  
6   guestroom            545 non-null   object  
7   basement              545 non-null   object  
8   hotwaterheating      545 non-null   object  
9   airconditioning      545 non-null   object  
10  parking              545 non-null   int64  
11  prefarea             545 non-null   object  
12  furnishingstatus     545 non-null   object  
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
```

```
In [5]: df.isnull().sum()
```

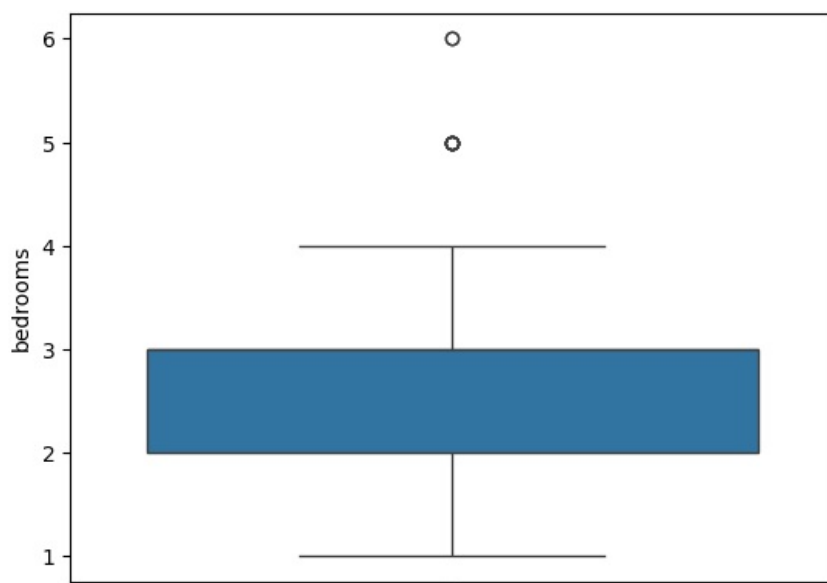
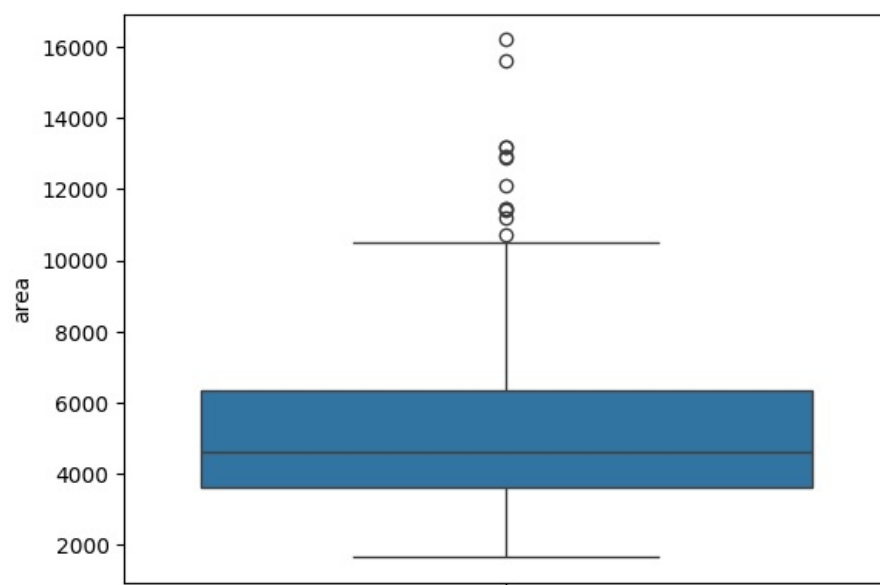
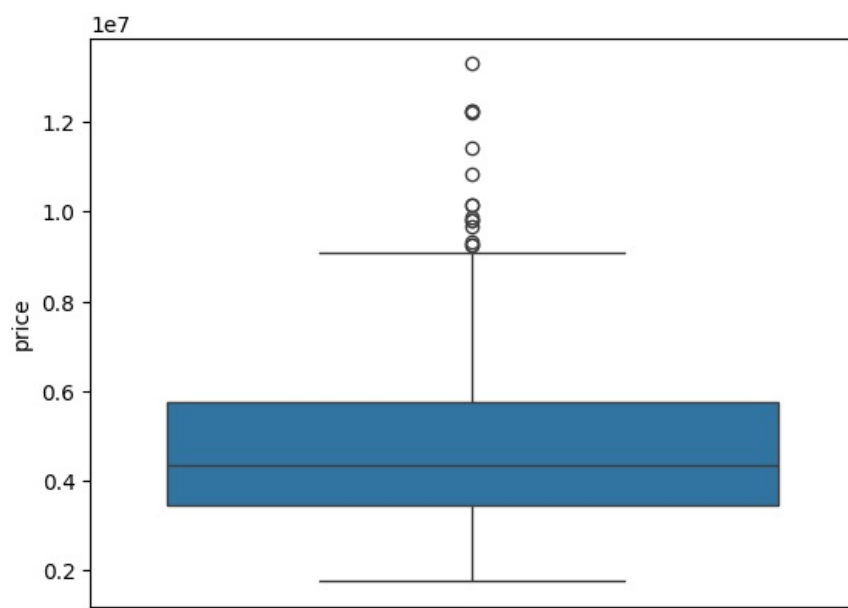
```
Out[5]: price                0
area                  0
bedrooms              0
bathrooms             0
stories               0
mainroad              0
guestroom             0
basement              0
hotwaterheating       0
airconditioning        0
parking               0
prefarea              0
furnishingstatus      0
dtype: int64
```

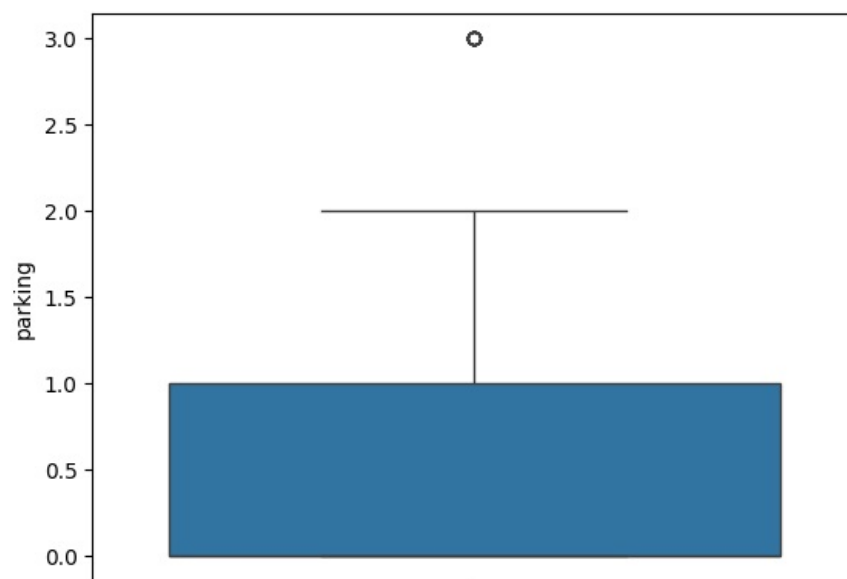
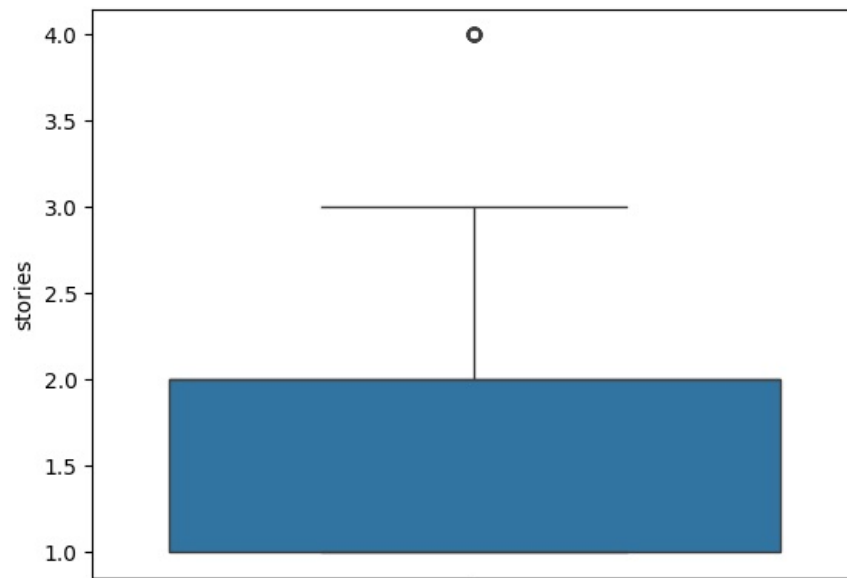
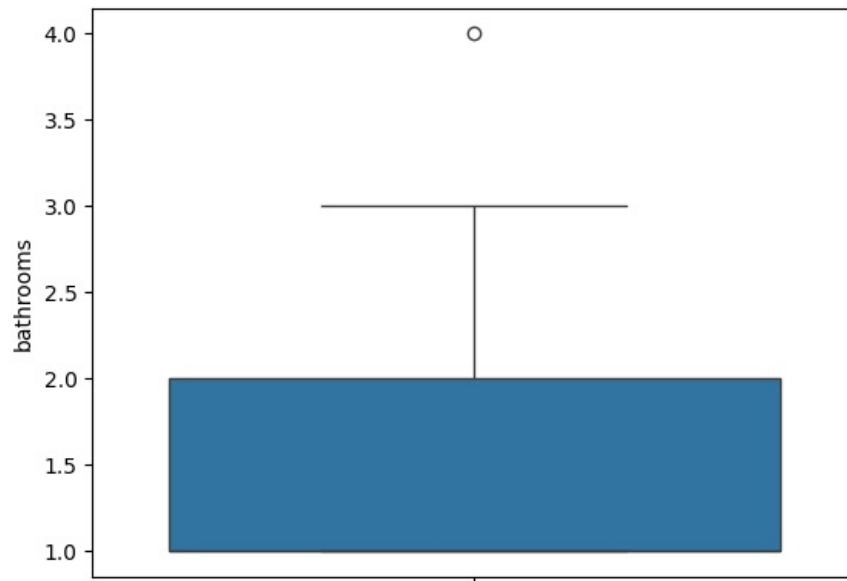
```
In [5]: # no null values in data
```

```
In [6]: #using boxplot to visualise any outliers presence
```

```
In [6]: df_numeric= df.select_dtypes(include='number')
```

```
In [7]: for i in df_numeric.columns:  
        sns.boxplot(df_numeric[i])  
        plt.show()
```



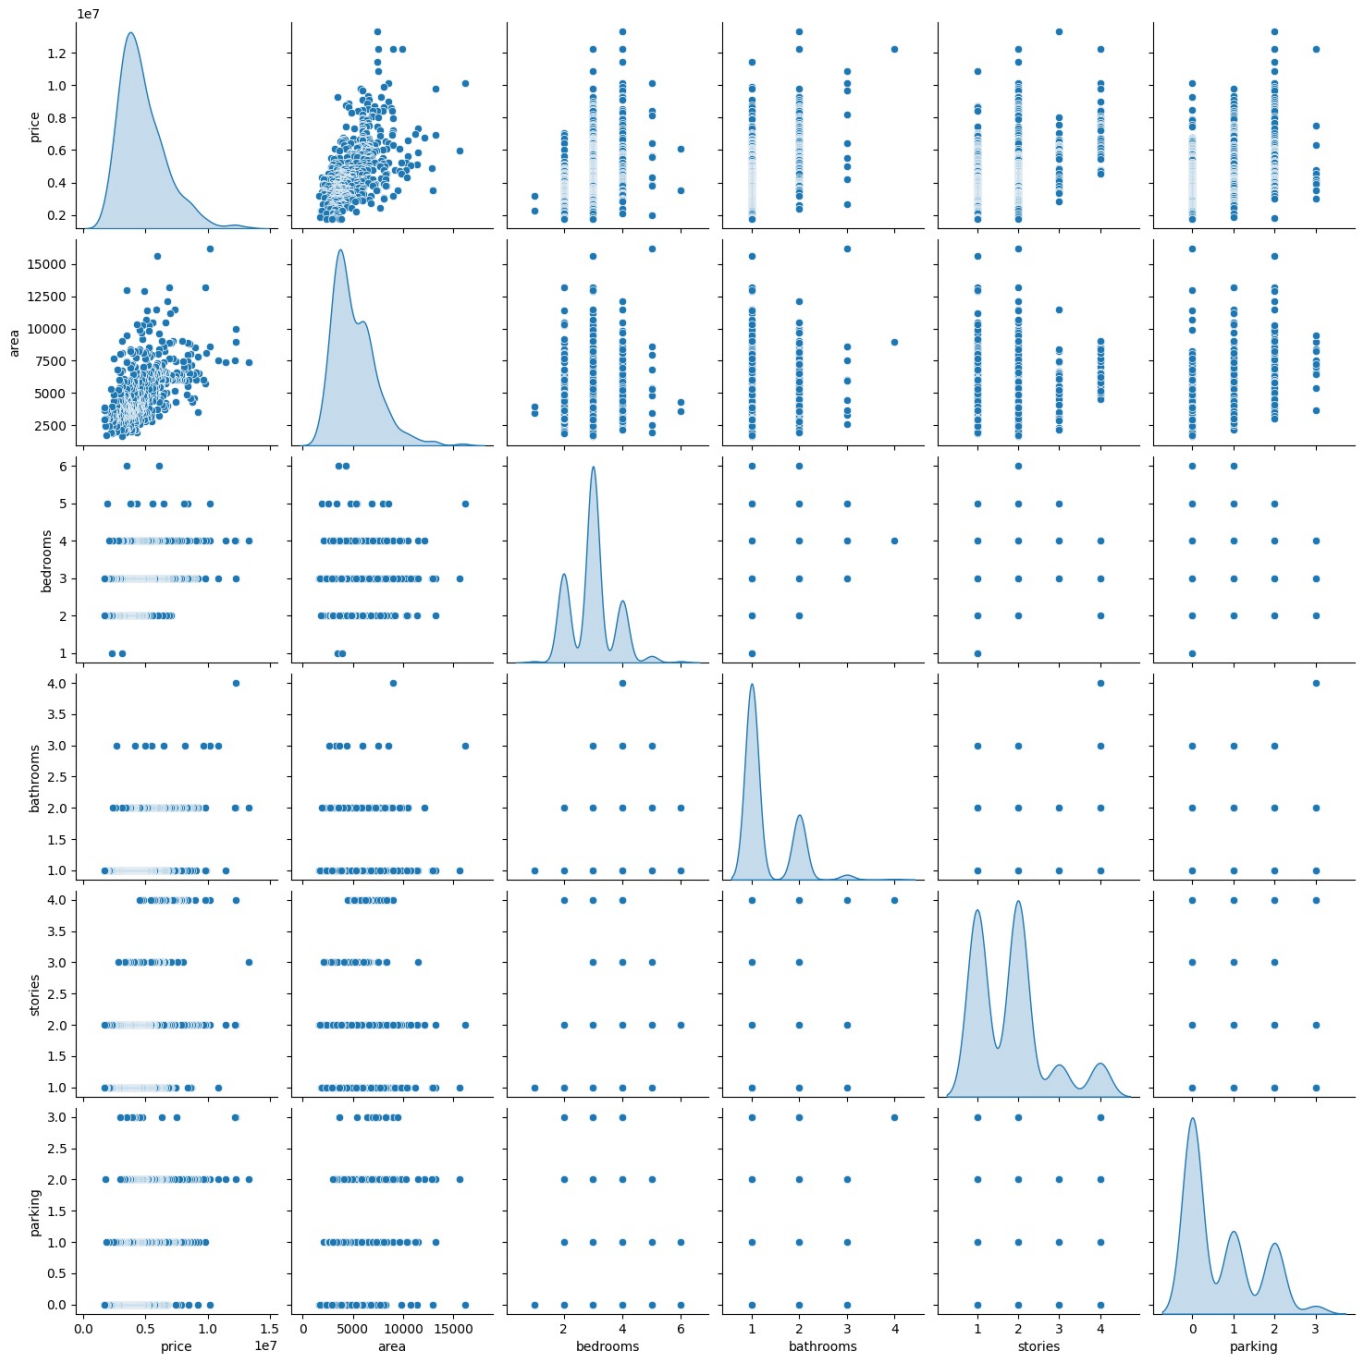


```
In [9]: # The presence of outliers is indicated
```

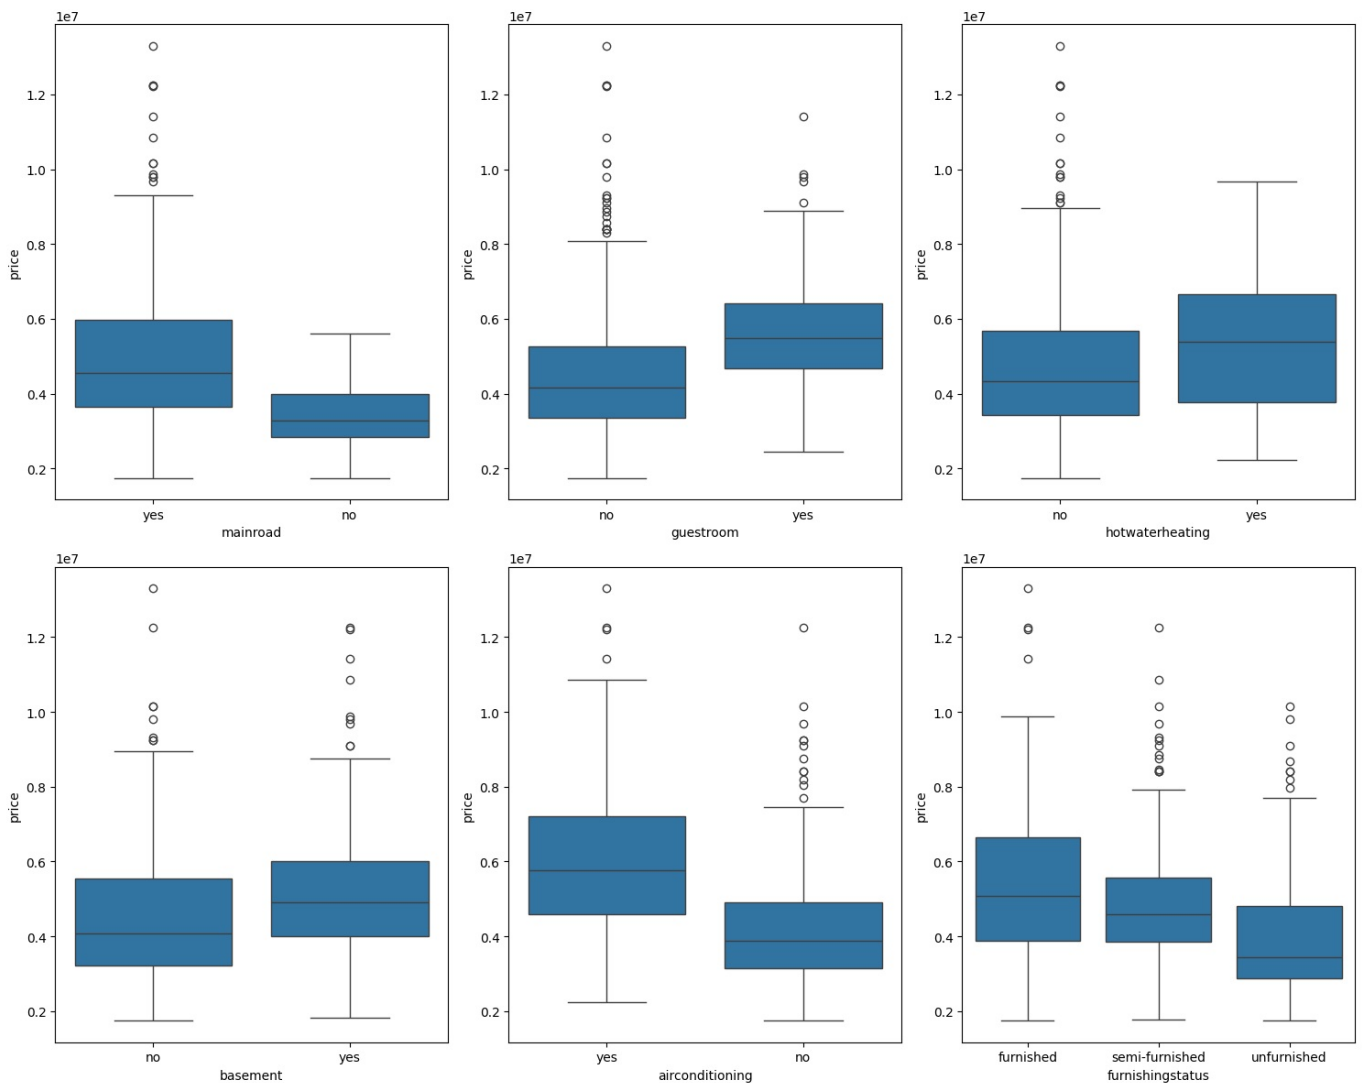
```
In [10]: # since the observations are very few.. we try not removing outliers
```

STEP3: visualising numeric variables

```
In [8]: sns.pairplot(df,diag_kind='kde')  
plt.show()
```



```
In [9]: # bivariate  
plt.figure(figsize=(15,12))  
plt.subplot(2,3,1)  
sns.boxplot(x='mainroad',y='price',data=df)  
plt.subplot(2,3,2)  
sns.boxplot(x='guestroom',y='price',data=df)  
plt.subplot(2,3,3)  
sns.boxplot(x='hotwaterheating',y='price',data=df)  
plt.subplot(2,3,4)  
sns.boxplot(x='basement',y='price',data=df)  
plt.subplot(2,3,5)  
sns.boxplot(x='airconditioning',y='price',data=df)  
plt.subplot(2,3,6)  
sns.boxplot(x='furnishingstatus',y='price',data=df)  
plt.tight_layout()  
plt.show()
```



```
In [14]: # if we look at the above boxplot the houses which are on mainroad, having guestroom, having hotwater facility,
# with basement, airconditioning and semi furnished and fully furnished are expensive and more than compared to
# houses with not having most facilities
```

```
In [10]: df.replace({'yes': 1, "no": 0},inplace=True)
# replacing object variables 'YES' and "NO" to one and zero
```

```
In [11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   price            545 non-null    int64
1   area             545 non-null    int64
2   bedrooms         545 non-null    int64
3   bathrooms        545 non-null    int64
4   stories          545 non-null    int64
5   mainroad         545 non-null    int64
6   guestroom        545 non-null    int64
7   basement         545 non-null    int64
8   hotwaterheating  545 non-null    int64
9   airconditioning  545 non-null    int64
10  parking          545 non-null    int64
11  prefarea         545 non-null    int64
12  furnishingstatus 545 non-null    object
dtypes: int64(12), object(1)
memory usage: 55.5+ KB
```

```
In [17]: # encoding furnishing status
```

```
In [12]: encoded_furniture=pd.get_dummies(df['furnishingstatus'],drop_first=True).astype('uint8')
```

```
In [13]: df= pd.concat([df, encoded_furniture], axis = 1)
df
```

Out[13]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	p
0	13300000	7420	4	2	3	1	0	0	0	1	2	
1	12250000	8960	4	4	4	1	0	0	0	1	3	
2	12250000	9960	3	2	2	1	0	1	0	0	2	
3	12215000	7500	4	2	2	1	0	1	0	1	3	
4	11410000	7420	4	1	2	1	1	1	0	1	2	
...
540	1820000	3000	2	1	1	1	0	1	0	0	2	
541	1767150	2400	3	1	1	0	0	0	0	0	0	
542	1750000	3620	2	1	1	1	0	0	0	0	0	
543	1750000	2910	3	1	1	0	0	0	0	0	0	
544	1750000	3850	3	1	2	1	0	0	0	0	0	

545 rows × 15 columns



In [14]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   price               545 non-null   int64
1   area                545 non-null   int64
2   bedrooms            545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   int64
6   guestroom            545 non-null   int64
7   basement             545 non-null   int64
8   hotwaterheating      545 non-null   int64
9   airconditioning      545 non-null   int64
10  parking              545 non-null   int64
11  prefarea             545 non-null   int64
12  furnishingstatus     545 non-null   object
13  semi-furnished       545 non-null   uint8
14  unfurnished          545 non-null   uint8
dtypes: int64(12), object(1), uint8(2)
memory usage: 56.5+ KB
```

In [21]: # we can drop furnishingstatus now

In [15]: df.drop('furnishingstatus',axis=1,inplace=True)

In [16]: # standardising the data
from sklearn.preprocessing import MinMaxScaler

In [17]: mm=MinMaxScaler()
df_scaled=pd.DataFrame(mm.fit_transform(df),columns=df.columns)
df_scaled

Out[17]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	park
0	1.000000	0.396564	0.6	0.333333	0.666667	1.0	0.0	0.0	0.0	1.0	0.666
1	0.909091	0.502405	0.6	1.000000	1.000000	1.0	0.0	0.0	0.0	1.0	1.000
2	0.909091	0.571134	0.4	0.333333	0.333333	1.0	0.0	1.0	0.0	0.0	0.666
3	0.906061	0.402062	0.6	0.333333	0.333333	1.0	0.0	1.0	0.0	1.0	1.000
4	0.836364	0.396564	0.6	0.000000	0.333333	1.0	1.0	1.0	0.0	1.0	0.666
...
540	0.006061	0.092784	0.2	0.000000	0.000000	1.0	0.0	1.0	0.0	0.0	0.666
541	0.001485	0.051546	0.4	0.000000	0.000000	0.0	0.0	0.0	0.0	0.0	0.000
542	0.000000	0.135395	0.2	0.000000	0.000000	1.0	0.0	0.0	0.0	0.0	0.000
543	0.000000	0.086598	0.4	0.000000	0.000000	0.0	0.0	0.0	0.0	0.0	0.000
544	0.000000	0.151203	0.4	0.000000	0.333333	1.0	0.0	0.0	0.0	0.0	0.000

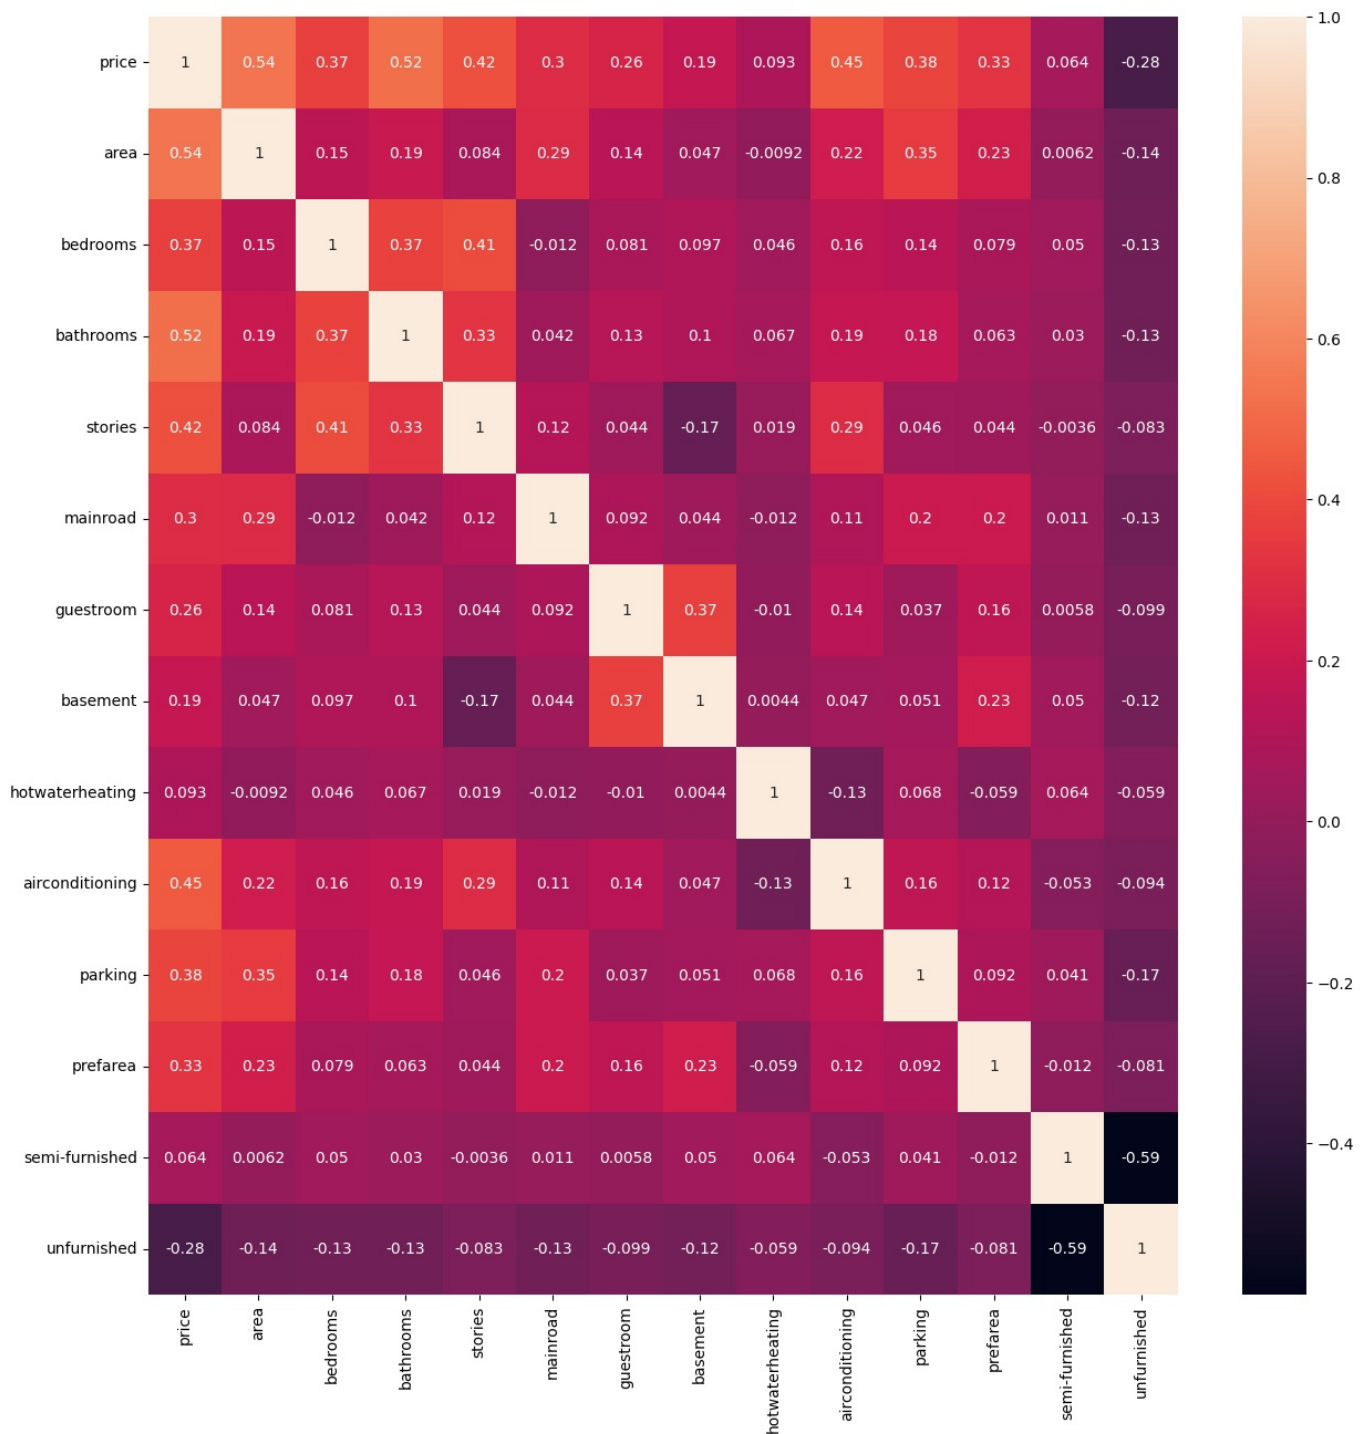
545 rows × 14 columns



In [18]:

```
plt.figure(figsize=(15,15))
sns.heatmap(df_scaled.corr(),annot=True)
```

Out[18]: <Axes: >



```
In [26]: # Area,bathrooms,bedrooms,stories with price look to have a good correlation,
```

STEP4: Splitting the Data into Training and Testing Sets

```
In [19]: from sklearn.model_selection import train_test_split
```

```
In [20]: X=df_scaled.drop('price',axis=1)
y=df_scaled['price']
```

```
In [21]: xtrain,xtest,ytrain,ytest= train_test_split(X,y, test_size = 0.2, random_state = 100)
xtrain.shape,xtest.shape,ytrain.shape,ytest.shape
```

```
Out[21]: ((436, 13), (109, 13), (436,), (109,))
```

STEP5: Model Building

```
In [22]: from sklearn.linear_model import LinearRegression
```

```
In [26]: lin_reg = LinearRegression()
# train the model with input and output data - train
model = lin_reg.fit(xtrain,ytrain)
```



```
#test the model with input data - test
y_pred = model.predict(xtest)
lin_reg.fit(xtest,y_pred)
```

Out[26]:

```
▼ LinearRegression ⓘ ?
LinearRegression()
```

In [32]: *#Model Evaluation*

In [27]: **from** sklearn.metrics **import** mean_squared_error,r2_score

In [34]: *# model evaluation for training set*
y_pred = lin_reg.predict(xtrain)
rmse = (np.sqrt(mean_squared_error(ytrain, y_pred)))
print("rmse:",rmse)
R square value - training dataset
r2 = r2_score(ytrain, y_pred)
print("R- Square:",r2)

rmse: 0.09160650416956634
R- Square: 0.6781789392869301

In [35]: *# model evaluation for testing set*
y_pred = lin_reg.predict(xtest)
rmse = (np.sqrt(mean_squared_error(ytest, y_pred)))
print("rmse:",rmse)
r2 = r2_score(ytest, y_pred)
print("R- Square:",r2)

rmse: 0.09191877175078576
R- Square: 0.6806539407870682

OLS - model summary

In [36]: **import** statsmodels.api **as** sm
model = sm.OLS(y, X).fit()
predictions = model.predict(X) *# make the predictions by the model*

Print out the statistics
model.summary()

Out[36]:

OLS Regression Results

Dep. Variable:	price	R-squared (uncentered):	0.912
Model:	OLS	Adj. R-squared (uncentered):	0.909
Method:	Least Squares	F-statistic:	421.6
Date:	Mon, 27 Nov 2023	Prob (F-statistic):	3.45e-270
Time:	15:07:26	Log-Likelihood:	530.67
No. Observations:	545	AIC:	-1035.
Df Residuals:	532	BIC:	-979.4
Df Model:	13		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
area	0.3137	0.030	10.399	0.000	0.254	0.373
bedrooms	0.0707	0.026	2.727	0.007	0.020	0.122
bathrooms	0.2531	0.027	9.480	0.000	0.201	0.306
stories	0.1161	0.017	6.974	0.000	0.083	0.149
mainroad	0.0443	0.010	4.260	0.000	0.024	0.065
guestroom	0.0260	0.011	2.281	0.023	0.004	0.048
basement	0.0315	0.009	3.318	0.001	0.013	0.050
hotwaterheating	0.0751	0.019	3.890	0.000	0.037	0.113
airconditioning	0.0756	0.009	8.077	0.000	0.057	0.094
parking	0.0718	0.015	4.721	0.000	0.042	0.102
prefarea	0.0558	0.010	5.574	0.000	0.036	0.075
semi-furnished	0.0009	0.009	0.100	0.921	-0.017	0.019
unfurnished	-0.0292	0.009	-3.075	0.002	-0.048	-0.011

Omnibus:	90.675	Durbin-Watson:	1.240
Prob(Omnibus):	0.000	Jarque-Bera (JB):	232.079
Skew:	0.841	Prob(JB):	4.02e-51
Kurtosis:	5.718	Cond. No.	11.1

Notes:

[1] R² is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Building random forest model with all features to find out the feature importance

In [37]:

```
from sklearn.ensemble import RandomForestRegressor

forest = RandomForestRegressor(n_estimators = 100, max_depth = 10)
forest.fit(xtrain,ytrain)
```

Out[37]:

RandomForestRegressor

RandomForestRegressor(max_depth=10)

In [38]:

```
y_pred=forest.predict(xtest)
```

In [39]:

```
rmse=np.sqrt(mean_squared_error(ytest,y_pred))
rmse
```

Out[39]:

0.0963470346104841

In [40]:

```
r2_score(ytest,y_pred)
```

Out[40]:

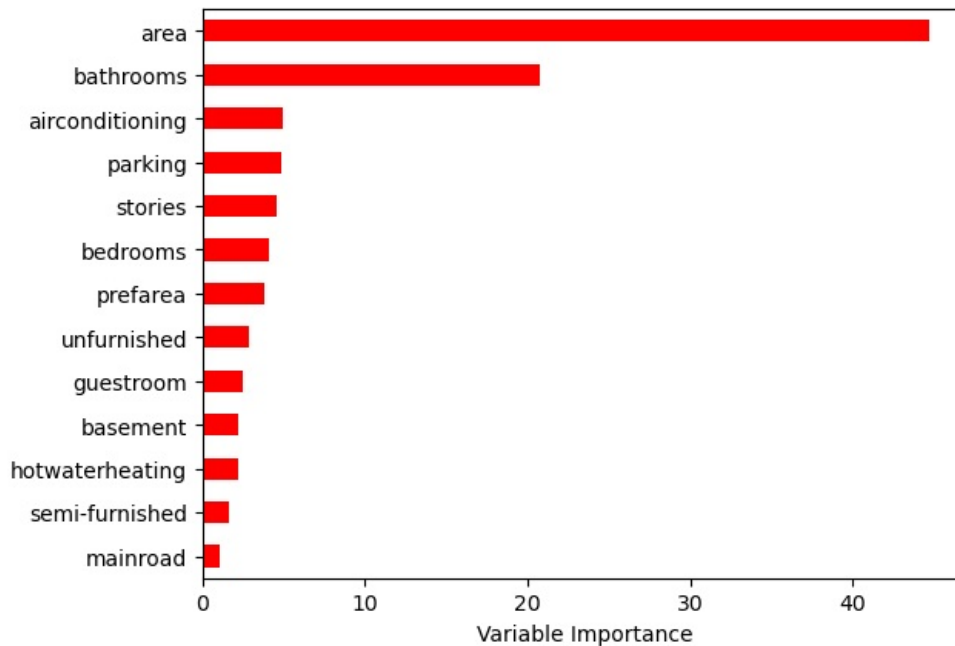
0.6491432461227882

In [41]:

```
Importance = pd.DataFrame({'Importance':forest.feature_importances_*100}, index=xtrain.columns)
Importance.sort_values('Importance', axis=0, ascending=True).plot(kind='barh', color='r', )
plt.xlabel('Variable Importance')
plt.gca().legend_ = None
```

```
plt.figure(figsize=(20,15))
```

Out[41]: <Figure size 2000x1500 with 0 Axes>



<Figure size 2000x1500 with 0 Axes>

```
In [42]: Importance.sort_values('Importance', axis=0, ascending=False)
```

Out[42]:

	Importance
area	44.750199
bathrooms	20.754096
airconditioning	4.918742
parking	4.817895
stories	4.517445
bedrooms	4.048446
prefarea	3.763245
unfurnished	2.891752
guestroom	2.485536
basement	2.240064
hotwaterheating	2.205662
semi-furnished	1.588988
mainroad	1.017932

```
In [43]: # area and bathrooms explain 65 percent variance together considering they are the most important features
```

```
In [44]: # creating new X based on selective features
# taking features that could give 85% variance together
X_new=df_scaled.drop(['unfurnished','guestroom','basement','hotwaterheating','semi-furnished','mainroad'],axis=1)
Y=df_scaled['price']
```

```
In [ ]:
```

```
In [48]: from sklearn import metrics
from sklearn.model_selection import KFold
from sklearn.linear_model import Ridge,Lasso
```

```
In [49]: LR=LinearRegression()
Ridge_R=Ridge(alpha=0.5)
Lasso_R=Lasso(alpha=0.1)
```

```
In [53]: kf=KFold(n_splits=3,shuffle=True,random_state=2)
for model, name in zip([LR,Ridge_R,Lasso_R],['Ridge','Lasso']):
    rmse=[]
    for train,test in kf.split(X_new,Y):
        x_train,x_test=X_new.iloc[train,:],X_new.iloc[test,:]
        y_train,y_test=Y.iloc[train],Y.iloc[test]
        model.fit(x_train,y_train)
        Y_predict=model.predict(x_test)
```

```
rmse.append(np.sqrt(metrics.mean_squared_error(y_test,Y_predict)))
print(rmse)
print("Cross_Validated_rmse_score: %0.03f (+/- %0.5f) [%s]" % (np.mean(rmse),np.var(rmse,ddof=1),name))
```

```
[1.2334373077647615e-16, 1.661481330599047e-16, 8.773567958375563e-17]
Cross_Validated_rmse_score: 0.000 (+/- 0.00000) [Ridge]
[0.012877039040894732, 0.011901611855598371, 0.012292062950779137]
Cross_Validated_rmse_score: 0.012 (+/- 0.00000) [Lasso]
```

Model Evaluation Summary

Initial Model: Linear Regression with All Features (Scaled) Root Mean Square Error (RMSE): 0.0916 Intermediate Model: Random Forest with All Features RMSE: 0.0958 Final Model: Linear Regression with Ridge and Lasso Regression (Selected Features)

Features selected to account for 85% variance

Ridge Regression:

Cross-Validated RMSE Score: 0.012 (\pm 0.00000)

Lasso Regression:

Cross-Validated RMSE Score: 0.00 (\pm 0.00000)

Conclusion

The model initially had an RMSE of 0.0916. A Random Forest model yielded a slightly higher RMSE of 0.0958. Using Ridge and Lasso regression with selected features improved the RMSE significantly. The Cross-Validated RMSE for Ridge Regression is 0.012, while for Lasso Regression it is 0.00. This demonstrates a significant reduction in RMSE, indicating that Ridge and Lasso regression models are more effective.

Measures of a Good Model

Lower RMSE: Indicates better predictive accuracy. Cross-Validation: Ensures the model's reliability and generalizability. Feature Importance: Identifying key features that contribute most to the prediction, improving model interpretability and performance. Variance Explained: Selecting features that explain a high percentage of variance helps in reducing complexity and enhancing model performance.

In []:

In []:

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js